

Heart Disease Risk Analysis

March 19, 2025

1 Heart Disease Risk Analysis Using EDA and Machine Learning

1.1 Step 1: Understanding the Data

```
[10]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from scipy.stats import ttest_ind
from scipy.stats import chi2_contingency
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, \
    f1_score, classification_report
from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, \
    f1_score
from sklearn.model_selection import cross_val_score
import shap
```

```
[11]: #data loading
df = pd.read_csv(r'C:\Users\User\Downloads\HeartDiseaseTrain-Test.csv')
```

```
[14]: #exploring the dataset
df.head()
```

```
[14]:   age    sex chest_pain_type  resting_blood_pressure  cholestoral  \
0   52  Male  Typical angina                125             212
1   53  Male  Typical angina                140             203
2   70  Male  Typical angina                145             174
3   61  Male  Typical angina                148             203
4   62 Female  Typical angina                138             294
```

	fasting_blood_sugar	rest_ecg	Max_heart_rate	\
0	Lower than 120 mg/ml	ST-T wave abnormality	168	
1	Greater than 120 mg/ml	Normal	155	
2	Lower than 120 mg/ml	ST-T wave abnormality	125	
3	Lower than 120 mg/ml	ST-T wave abnormality	161	
4	Greater than 120 mg/ml	ST-T wave abnormality	106	

	exercise_induced_angina	oldpeak	slope	vessels_colored_by_flourosopy	\
0	No	1.0	Downsloping	Two	
1	Yes	3.1	Upsloping	Zero	
2	Yes	2.6	Upsloping	Zero	
3	No	0.0	Downsloping	One	
4	No	1.9	Flat	Three	

	thalassemia	target
0	Reversible Defect	0
1	Reversible Defect	0
2	Reversible Defect	0
3	Reversible Defect	0
4	Fixed Defect	0

```
[16]: df.rename(columns={'cholestorol': 'cholesterol'}, inplace=True)
```

```
[18]: # datatypes
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                    1025 non-null   int64
1   sex                                    1025 non-null   object
2   chest_pain_type                       1025 non-null   object
3   resting_blood_pressure                1025 non-null   int64
4   cholesterol                           1025 non-null   int64
5   fasting_blood_sugar                   1025 non-null   object
6   rest_ecg                             1025 non-null   object
7   Max_heart_rate                        1025 non-null   int64
8   exercise_induced_angina              1025 non-null   object
9   oldpeak                              1025 non-null   float64
10  slope                                1025 non-null   object
11  vessels_colored_by_flourosopy         1025 non-null   object
12  thalassemia                           1025 non-null   object
13  target                                1025 non-null   int64
dtypes: float64(1), int64(5), object(8)
memory usage: 112.2+ KB
```

```
[20]: # Summary statistics for numeric columns
df.describe()
```

```
[20]:
```

	age	resting_blood_pressure	cholesterol	Max_heart_rate \
count	1025.000000	1025.000000	1025.00000	1025.000000
mean	54.434146	131.611707	246.00000	149.114146
std	9.072290	17.516718	51.59251	23.005724
min	29.000000	94.000000	126.00000	71.000000
25%	48.000000	120.000000	211.00000	132.000000
50%	56.000000	130.000000	240.00000	152.000000
75%	61.000000	140.000000	275.00000	166.000000
max	77.000000	200.000000	564.00000	202.000000

	oldpeak	target
count	1025.000000	1025.000000
mean	1.071512	0.513171
std	1.175053	0.500070
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.800000	1.000000
75%	1.800000	1.000000
max	6.200000	1.000000

```
[22]: #identify missing data
df.isnull().sum()
```

```
[22]: age                                0
sex                                      0
chest_pain_type                        0
resting_blood_pressure                 0
cholesterol                           0
fasting_blood_sugar                   0
rest_ecg                             0
Max_heart_rate                        0
exercise_induced_angina               0
oldpeak                              0
slope                                0
vessels_colored_by_flourosopy         0
thalassemia                          0
target                               0
dtype: int64
```

```
[24]: #identify duplicates & Check number of duplicates
df.duplicated().sum()
```

```
[24]: 723
```

```
[26]: # Identify categorical columns
categorical_cols = df.select_dtypes(include=['object', 'category']).columns.
      ↪tolist()
print("Categorical Columns:", categorical_cols)
```

Categorical Columns: ['sex', 'chest_pain_type', 'fasting_blood_sugar', 'rest_ecg', 'exercise_induced_angina', 'slope', 'vessels_colored_by_flourosopy', 'thalassemia']

```
[28]: # Check which columns exist in the dataset
valid_categorical_cols = [col for col in categorical_cols if col in df.columns]
print("Valid columns for Label Encoding:", valid_categorical_cols)
```

Valid columns for Label Encoding: ['sex', 'chest_pain_type', 'fasting_blood_sugar', 'rest_ecg', 'exercise_induced_angina', 'slope', 'vessels_colored_by_flourosopy', 'thalassemia']

```
[30]: # Apply Label Encoding to existing categorical columns
label_encoder = LabelEncoder()
for col in valid_categorical_cols:
    df[col] = label_encoder.fit_transform(df[col])
```

```
[32]: # Verify encoding
print(df.head())
```

	age	sex	chest_pain_type	resting_blood_pressure	cholesterol	\
0	52	1	3	125	212	
1	53	1	3	140	203	
2	70	1	3	145	174	
3	61	1	3	148	203	
4	62	0	3	138	294	

	fasting_blood_sugar	rest_ecg	Max_heart_rate	exercise_induced_angina	\
0	1	2	168	0	
1	0	1	155	1	
2	1	2	125	1	
3	1	2	161	0	
4	0	2	106	0	

	oldpeak	slope	vessels_colored_by_flourosopy	thalassemia	target
0	1.0	0	3	3	0
1	3.1	2	4	3	0
2	2.6	2	4	3	0
3	0.0	0	1	3	0
4	1.9	1	2	0	0

2 Normalization

```
[35]: # Identify numerical columns (integers & floats) excluding the target
numerical_cols = df.select_dtypes(include=['int64', 'float64']).columns.tolist()
numerical_cols.remove('target')
```

```
[37]: # Apply Min-Max Normalization
scaler = MinMaxScaler()
df_normalized = df.copy()
df_normalized[numerical_cols] = scaler.fit_transform(df[numerical_cols])
```

```
[39]: # Apply Min-Max Normalization
scaler = MinMaxScaler()
df_normalized = df.copy()
df_normalized[numerical_cols] = scaler.fit_transform(df[numerical_cols])
```

```
[41]: # Display results
print(df_normalized.head()) # Normalized dataset
```

	age	sex	chest_pain_type	resting_blood_pressure	cholesterol	\
0	0.479167	1	3	0.292453	0.196347	
1	0.500000	1	3	0.433962	0.175799	
2	0.854167	1	3	0.481132	0.109589	
3	0.666667	1	3	0.509434	0.175799	
4	0.687500	0	3	0.415094	0.383562	

	fasting_blood_sugar	rest_ecg	Max_heart_rate	exercise_induced_angina	\
0		1	2	0.740458	0
1		0	1	0.641221	1
2		1	2	0.412214	1
3		1	2	0.687023	0
4		0	2	0.267176	0

	oldpeak	slope	vessels_colored_by_flourosopy	thalassemia	target
0	0.161290	0	3	3	0
1	0.500000	2	4	3	0
2	0.419355	2	4	3	0
3	0.000000	0	1	3	0
4	0.306452	1	2	0	0

3 EXPLORATORY ANALYSIS

```
[44]: # Define numerical columns first
num_cols = ['age', 'cholesterol', 'resting_blood_pressure', 'Max_heart_rate']

# Generates unique colors
palette = sns.color_palette("husl", len(num_cols))
```

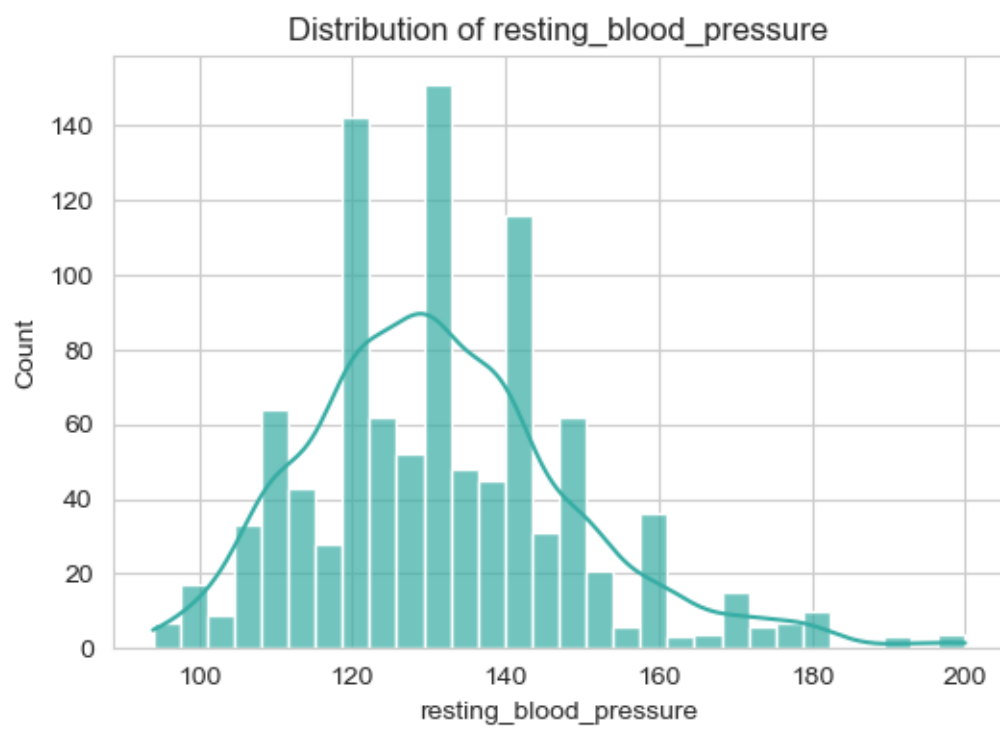
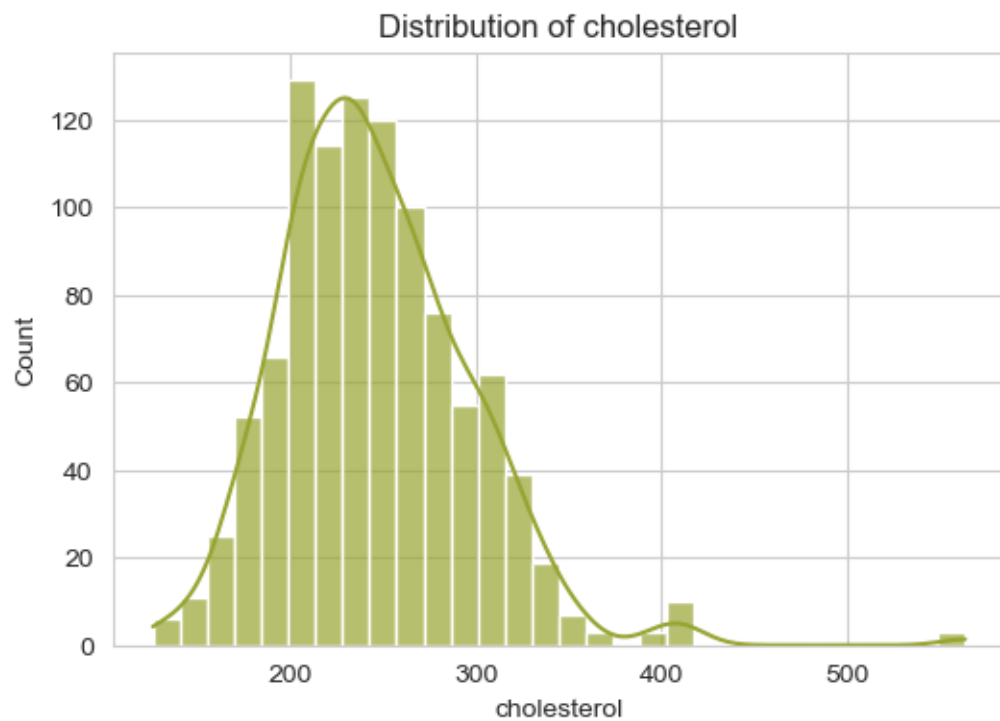
```

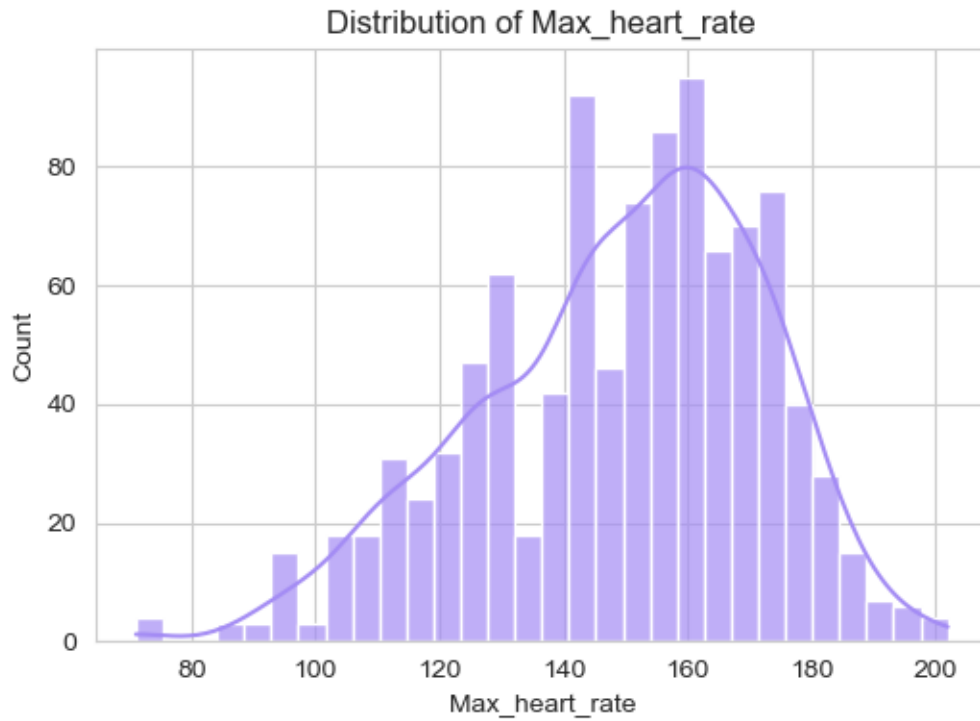
# Set a nice background
sns.set_style("whitegrid")

# Loop through numerical columns and plot distributions
for col, color in zip(num_cols, palette):
    if col in df.columns:
        plt.figure(figsize=(6, 4))
        sns.histplot(df[col], bins=30, kde=True, color=color, alpha=0.7)
        plt.title(f"Distribution of {col}")
        plt.show()

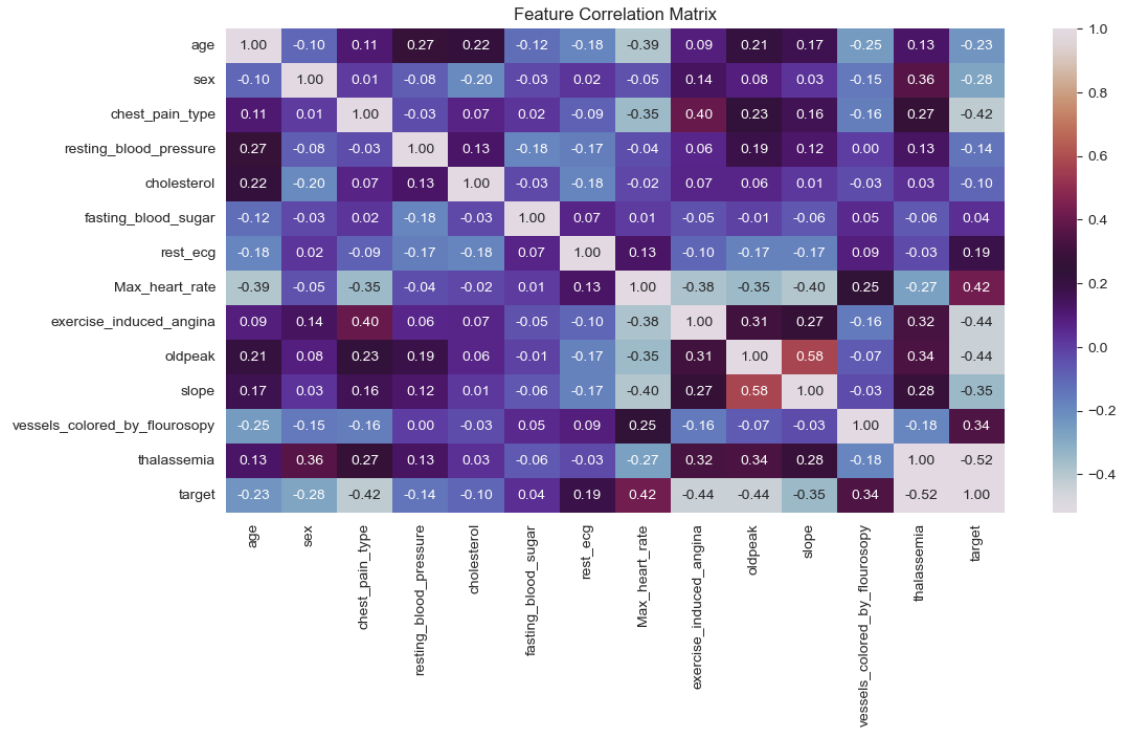
```



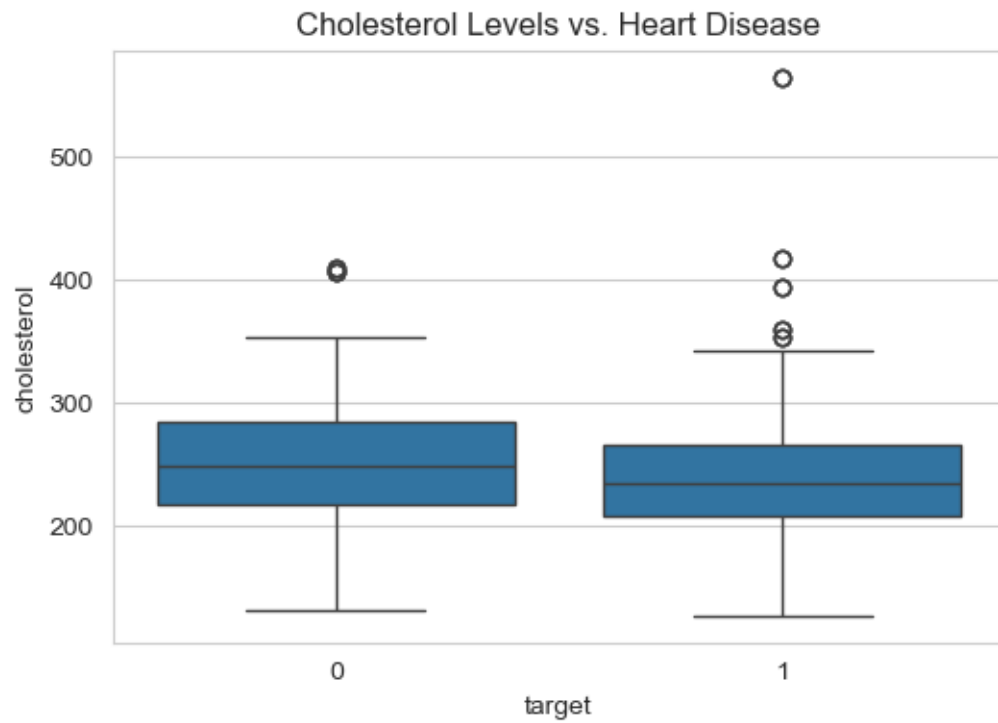




```
[46]: #Bivariate Analysis - Relationship with Heart Disease
plt.figure(figsize=(12, 6))
sns.heatmap(df.corr(), annot=True, cmap="twilight", fmt=".2f")
plt.title("Feature Correlation Matrix")
plt.show()
```

```
[48]: # Boxplot of cholesterol levels by heart disease status
plt.figure(figsize=(6, 4))
sns.boxplot(x=df['target'], y=df['cholesterol'])
plt.title("Cholesterol Levels vs. Heart Disease")
plt.show()
```



```
[50]: # Multivariate Analysis - Pair Plot
sns.pairplot(df, hue="target", vars=['age', 'cholesterol', 'resting_blood_pressure', 'Max_heart_rate'])
plt.show()
```



4 Hypothesis Testing

5 Test if Cholesterol Levels Differ Between Heart Disease and No Heart Disease Patients

5.1 Hypothesis:

5.2 Null (H₀): No difference in cholesterol levels between patients with and without heart disease.

5.3 Alternative (H_a): Cholesterol levels significantly differ between the two groups.

```
[53]: # Step 1: Remove duplicate column (keep only one 'cholesterol')
df = df.loc[:, ~df.columns.duplicated()]

# Step 2: Ensure 'cholesterol' is a single column
df['cholesterol'] = pd.to_numeric(df['cholesterol'], errors='coerce')

print("Successfully fixed 'cholesterol' column and converted to numeric.")
```

Successfully fixed 'cholesterol' column and converted to numeric.

```
[56]: # Grouping by heart disease presence
heart_disease = df[df['target'] == 1]['cholesterol']
no_heart_disease = df[df['target'] == 0]['cholesterol']

# Perform t-test
t_stat, p_value = ttest_ind(heart_disease, no_heart_disease, equal_var=False)
print(f"T-statistic: {t_stat}, P-value: {p_value}")

if p_value < 0.05:
    print("Cholesterol levels are significantly different between groups.")
else:
    print("No significant difference in cholesterol levels.")
```

T-statistic: -3.219133918272536, P-value: 0.0013262409956772104

Cholesterol levels are significantly different between groups.

6 Test if Sex is Associated with Heart Disease (Chi-Square Test)

6.1 Hypothesis:

6.2 H : Sex is not related to heart disease.

6.3 H : Sex is significantly associated with heart disease.

```
[59]: # Creating a contingency table
contingency_table = pd.crosstab(df['sex'], df['target'])

# Perform Chi-Square Test
chi2, p_value, _, _ = chi2_contingency(contingency_table)
print(f"Chi-Square Statistic: {chi2}, P-value: {p_value}")

if p_value < 0.05:
    print("Sex and heart disease are significantly associated.")
else:
    print("No significant association between sex and heart disease.")
```

Chi-Square Statistic: 78.86305133922973, P-value: 6.656820681726434e-19
Sex and heart disease are significantly associated.

7 Test if Blood Pressure is Higher in Patients with Heart Disease

7.1 Hypothesis:

7.2 H : No difference in resting blood pressure between groups.

7.3 H : Patients with heart disease have significantly higher blood pressure

```
[62]: # Grouping by heart disease presence
bp_heart_disease = df[df['target'] == 1]['resting_blood_pressure']
bp_no_heart_disease = df[df['target'] == 0]['resting_blood_pressure']

# Perform t-test
t_stat, p_value = ttest_ind(bp_heart_disease, bp_no_heart_disease,
                             equal_var=False)
print(f"T-statistic: {t_stat}, P-value: {p_value}")

if p_value < 0.05:
    print("Blood pressure is significantly different in heart disease patients.")
else:
    print("No significant difference in blood pressure levels.")
```

T-statistic: -4.465214972380933, P-value: 8.922491860767991e-06
Blood pressure is significantly different in heart disease patients.

8 Model Building

```
[65]: #split data before training
X = df.drop(columns=['target']) # Features
y = df['target'] # Target variable

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[67]: #Standardize data (important for KNN & Logistic Regression)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

8.1 Train Models

```
[70]: # Initialize models
models = {
    "Logistic Regression": LogisticRegression(),
    "Decision Tree": DecisionTreeClassifier(),
    "K-Nearest Neighbors": KNeighborsClassifier(n_neighbors=5)
}
```

```
[72]: #Train & Evaluate Models
for name, model in models.items():
    model.fit(X_train, y_train) # Train
    y_pred = model.predict(X_test) # Predict
    # Calculate metrics
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    print(f" {name} Performance:")
    print(f"   - Accuracy: {accuracy:.4f}")
    print(f"   - Precision: {precision:.4f}")
    print(f"   - Recall: {recall:.4f}")
    print(f"   - F1-score: {f1:.4f}")
    print("-" * 40)
```

Logistic Regression Performance:

- Accuracy: 0.7951
- Precision: 0.7699
- Recall: 0.8447
- F1-score: 0.8056

Decision Tree Performance:

- Accuracy: 0.9854
- Precision: 1.0000

- Recall: 0.9709
- F1-score: 0.9852

K-Nearest Neighbors Performance:

- Accuracy: 0.8244
 - Precision: 0.7965
 - Recall: 0.8738
 - F1-score: 0.8333
-

9 Hyperparameter Tuning

9.1 Decision Tree Hyperparameter Tuning

9.1.1 `max_depth` (controls the depth of the tree)

9.1.2 `min_samples_split` (minimum samples needed to split a node)

9.1.3 `min_samples_leaf` (minimum samples in a leaf)

```
[76]: # Define hyperparameter grids
dt_params = {
    'max_depth': [3, 5, 10, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 5]
}
```

9.2 KNN Hyperparameter Tuning

9.2.1 `n_neighbors` (number of nearest neighbors)

9.2.2 `weights` (uniform vs. distance)

9.2.3 `metric` (euclidean, manhattan)

```
[79]: knn_params = {
    'n_neighbors': [3, 5, 7, 9, 11],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan']
}
```

```
[81]: # Initialize models
dt = DecisionTreeClassifier(random_state=42)
knn = KNeighborsClassifier()
```

```
[83]: # Grid search with cross-validation
dt_grid = GridSearchCV(dt, dt_params, cv=5, scoring='accuracy', n_jobs=-1)
knn_grid = GridSearchCV(knn, knn_params, cv=5, scoring='accuracy', n_jobs=-1)
```

```
[85]: # Fit models
dt_grid.fit(X_train, y_train)
knn_grid.fit(X_train, y_train)
```

```
[85]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(), n_jobs=-1,
                  param_grid={'metric': ['euclidean', 'manhattan'],
                              'n_neighbors': [3, 5, 7, 9, 11],
                              'weights': ['uniform', 'distance']},
                  scoring='accuracy')
```

```
[86]: # Best models
best_dt = dt_grid.best_estimator_
best_knn = knn_grid.best_estimator_
```

```
[87]: # Print best parameters
print("Best Decision Tree Parameters:", dt_grid.best_params_)
print("Best KNN Parameters:", knn_grid.best_params_)
```

```
Best Decision Tree Parameters: {'max_depth': None, 'min_samples_leaf': 1,
'min_samples_split': 2}
```

```
Best KNN Parameters: {'metric': 'euclidean', 'n_neighbors': 5, 'weights':
'distance'}
```

9.3 Evaluate Tuned Models

let's train the optimized models and compare their accuracy, precision, recall, and F1-score.

```
[92]: # Train optimized Decision Tree
best_dt.fit(X_train, y_train)
y_pred_dt = best_dt.predict(X_test)
```

```
[94]: # Train optimized KNN
best_knn.fit(X_train, y_train)
y_pred_knn = best_knn.predict(X_test)
```

```
[96]: # Function to evaluate model
def evaluate_model(y_true, y_pred, model_name):
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred)
    recall = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)

    print(f" {model_name} Performance:")
    print(f"   - Accuracy: {accuracy:.4f}")
    print(f"   - Precision: {precision:.4f}")
    print(f"   - Recall: {recall:.4f}")
    print(f"   - F1-score: {f1:.4f}")
    print("-" * 40)
```



```
[98]: # Evaluate models
evaluate_model(y_test, y_pred_dt, "Tuned Decision Tree")
evaluate_model(y_test, y_pred_knn, "Tuned KNN")
```

Tuned Decision Tree Performance:

- Accuracy: 0.9854
- Precision: 1.0000
- Recall: 0.9709
- F1-score: 0.9852

Tuned KNN Performance:

- Accuracy: 1.0000
 - Precision: 1.0000
 - Recall: 1.0000
 - F1-score: 1.0000
-

9.4 Cross-Validation

Instead of relying only on one train-test split, let's use Stratified K-Fold Cross-Validation to ensure stable results.

```
[101]: # 10-fold Cross-validation
cv_scores_dt = cross_val_score(best_dt, X, y, cv=10, scoring='accuracy')
cv_scores_knn = cross_val_score(best_knn, X, y, cv=10, scoring='accuracy')

print(f" Decision Tree Cross-Validation Accuracy: {cv_scores_dt.mean():.4f} ±
↪{cv_scores_dt.std():.4f}")
print(f" KNN Cross-Validation Accuracy: {cv_scores_knn.mean():.4f} ±
↪{cv_scores_knn.std():.4f}")
```

Decision Tree Cross-Validation Accuracy: 0.9912 ± 0.0134

KNN Cross-Validation Accuracy: 0.9942 ± 0.0117

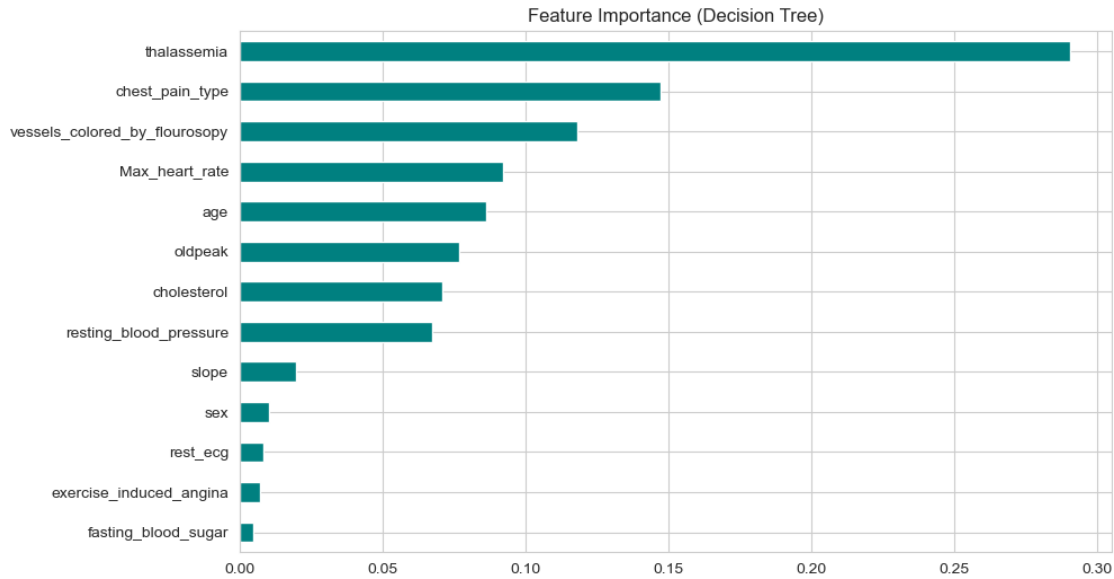
9.5 Check for Data Leakage

Feature Importance (Decision Tree)

If some features have extremely high importance, they might be leaking target information.

```
[104]: # Get feature importances
feat_importances = pd.Series(best_dt.feature_importances_, index=X.columns)
```

```
[106]: # Plot
plt.figure(figsize=(10, 6))
feat_importances.sort_values().plot(kind="barh", color="teal")
plt.title("Feature Importance (Decision Tree)")
plt.show()
```



Objective of SHAP Analysis ## Identify which features (like age, cholesterol, blood pressure, etc.) most influence the model's prediction. ## Determine whether features increase or decrease heart disease risk. ## Provide interpretable insights to make the model's decision-making transparent.

```
[109]: #Train a Machine Learning Model
#Before applying SHAP, ensure that the KNN classifier is properly trained:
knn.fit(X_train, y_train)
```

```
[109]: KNeighborsClassifier()
```

```
[111]: df.columns
```

```
[111]: Index(['age', 'sex', 'chest_pain_type', 'resting_blood_pressure',
        'cholesterol', 'fasting_blood_sugar', 'rest_ecg', 'Max_heart_rate',
        'exercise_induced_angina', 'oldpeak', 'slope',
        'vessels_colored_by_flourosopy', 'thalassemia', 'target'],
        dtype='object')
```

```
[199]: feature_names = [ 'cholesterol', 'age']
```

9.6 Define the SHAP Explainer

9.6.1 SHAP has different explainers, but since KNN doesn't have a direct SHAP integration, we use KernelExplainer:

```
[116]: # Select a small subset of training data for efficient computation
sample_data = X_train[:100]

# Define SHAP Explainer
explainer = shap.KernelExplainer(knn.predict_proba, sample_data)
```

```
[173]: # Generate SHAP values for a test sample
shap_values = explainer.shap_values(X_test)
```

```
0%|          | 0/205 [00:00<?, ?it/s]
```

```
[175]: # Check if SHAP values are in multi-class format (list of arrays)
if isinstance(shap_values, list) and len(shap_values) > 1:
    class_index = 1
    shap_values = shap_values[class_index]
# Print shape for verification
print("SHAP values shape after fix:", len(shap_values), len(shap_values[0]))
```

```
SHAP values shape after fix: 205 13
```

```
[177]: # Print shape for verification
print("SHAP values shape after fix:", len(shap_values), len(shap_values[0]))
```

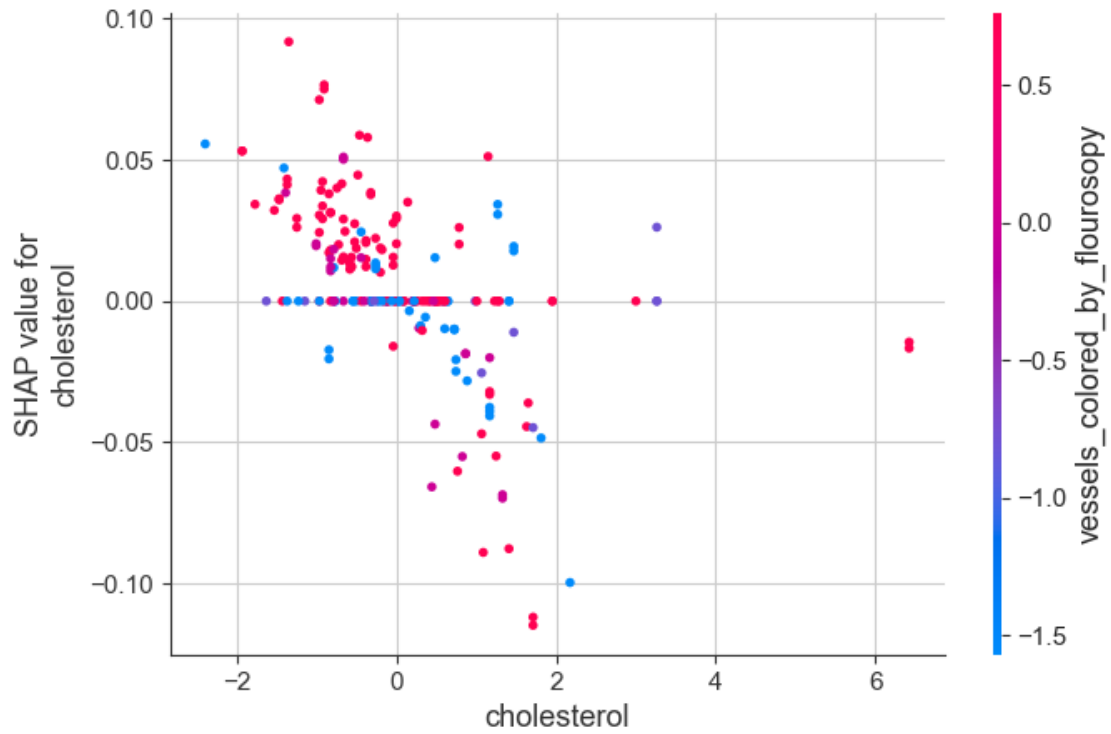
```
SHAP values shape after fix: 205 13
```

```
[179]: # Select SHAP values for one class
shap_values_class1 = shap_values[:, :, 1]

print("Fixed SHAP values shape:", shap_values_class1.shape)

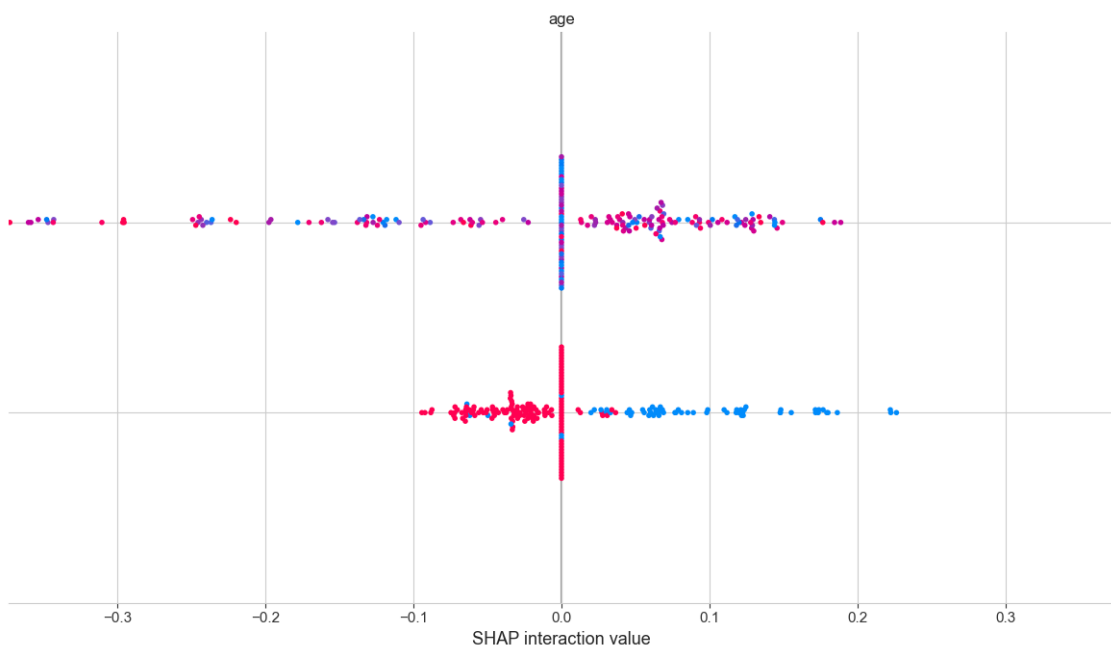
# plot
shap.dependence_plot("cholesterol", shap_values_class1, X_test,
    ↪feature_names=feature_names)
```

```
Fixed SHAP values shape: (205, 13)
```



```
[201]: shap.summary_plot(shap_values, X_test, feature_names=feature_names)
```

<Figure size 640x480 with 0 Axes>



10 Heart Disease Risk Analysis

10.1 The objective of the project was to analyze heart disease risk key factors using machine learning model. The dataset used is the UCI Heart Disease Dataset, which contains key health indicators such as cholesterol levels, blood pressure, and age. The workflow consists of the following steps.

10.1.1 1. Data Preprocessing: Handling missing values, feature scaling, and encoding categorical variables.

10.1.2 2. Exploratory Data Analysis (EDA): Understanding feature distributions, correlations, and key risk indicators.

10.1.3 3. Model Training & Evaluation: Training multiple machine learning models, selecting the best-performing one using metrics like accuracy, precision, and recall.

10.1.4 4. Feature Importance Analysis: Using SHAP to interpret the model's predictions and understand which features contribute most to heart disease risk.

10.1.5 5. Visualization & Insights: Generating SHAP plots and other visualizations to make the results interpretable.

This analysis helps in understanding the most critical factors influencing heart disease, providing a foundation for further healthcare-related predictive modeling and decision-making.

[]:

[]:

[]:

[]:

[]: